



[논문 리뷰] VGGNet (2014)

1. Introduction

☀ VGGNet은 16 ~ 19 weight layer이다. 해당 모델은 3x3 convolution filter를 사용함으로써, depths를 증가시켰습니다.

VGGNet 논문에서는 ConvNet architecture의 depth를 중요하게 다룹니다. depth를 다루기 위하여, Architecture의 parameter들을 변경하였고, **3x3 filter를 지속적으로 추가**함으로써, layer의 깊이를 증가시켰습니다.

☀ **결과적으로, 더 정확한 ConvNet구조를 구상하고, 다른 이미지 인식 dataset에도 적용가능하다.**

▼ 2. ConvNet Configurations

공정한 설정을 위해, ConvNet depth가 증가함에 따라, 정확한 evaluation을 위해,

ConvNet layer의 parameter를 동일하게 적용하였습니다.

2-1. Architecture

ConvNet에서는 (3, 224, 224) Image를 사용하였으며, 전처리 과정에서는 단지, **RGB value의 평균값을 training set의 각 pixel값 별로 subtract**해주었습니다.

Image는 반복적인 3x3 filter가 쌓인, stack of convolutional layer를 통과합니다. 또한, Input channels 에 비 선형성(linear transformation)을 적용하기 위해, 1x1 filter도 사용하였습니다.

대체적으로, stride=1을 적용하였으며, 공간 해상도(Spatial Resolution)를 보전하기 위해 padding을 적용하였습니다. 일부 Conv layer에는 Max-pooling(2x2, stride = 2)를 사용합니다.

Hidden layer에는 ReLU를 사용하였지만, AlexNet에서 사용한 Local Response Normalisation은 사용하지 않았습니다. 그러한 이유는, **memory 소비 및 computation time만 증가시키고, 성능향상에는 도움이 안되기 때문**입니다.

Conv layers이후에 3개의 FC가 존재하는데, 첫 번째, 두 번째는 4096 channels을 사용하고, 세 번째에서는 1000개의 channels을 사용하였습니다. 마지막 layer에서는 softmax를 사용하였습니다.

2-2. Configurations

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: **Number of parameters** (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

2-3. Discussion

VGGNet에서는 3x3 filter 여러개를 전체 NeuralNet에 걸쳐 사용합니다.

3x3 filter 2개를 사용하면, 5x5의 receptive field의 성능을 가지고, 3개를 사용하면 7x7 receptive field를 가집니다.

For instance, a stack of three 3x3 conv. layers instead of single 7x7 layer?

First, 단일 layer를 사용하는 것보다, 3개의 non-linear을 사용함으로써, 더욱 차별화된 의사결정을 하는 것이 가능합니다.

Second, Parameters의 수가 감소합니다.

Second, we decrease the number of parameters: assuming that both the input and the output of a three-layer 3×3 convolution stack has C channels, the stack is parametrised by $3(3^2 C^2) = 27C^2$ weights; at the same time, a single 7×7 conv. layer would require $7^2 C^2 = 49C^2$ parameters, i.e. 81% more. This can be seen as imposing a regularisation on the 7×7 conv. filters, forcing them to have a decomposition through the 3×3 filters (with non-linearity injected in between).

▼ 3. Classification Framework

3-1. Training

Initialisation of the network weights는 매우 중요합니다.

Random한 initialize로 훈련이 가능하게끔 하였습니다. 4개의 ConvNet과 3개의 FC를 사용하여, 학습을 시작합니다. 중간 layer들의 값을 $0 \sim 10^{-2}$ 의 가중치 값으로 초기화하였다.

224*224의 이미지를 사용하기 위해서, training image내에서 무작위로 cropped 하였다. 더 나아가, random horizontal flipping 과 random RGB colour shift를 사용하였습니다.

▼ 4. Classification Experiments

First, VGGNet에서 local response normalisation(A-LRN network)은 모든 normalisation layer가 없는 model에서 개선되지 않습니다. 그러므로, B ~ E 에서는 정규화를 사용하지 않습니다.

Second, ConvNet 깊이가 증가함에 따라, Classification Error가 감소하는 것을 알 수 있습니다.

✨ Scale jittering Augmentation에 의한 training은 Multi-scale image를 capturing하는데 도움을 줍니다.

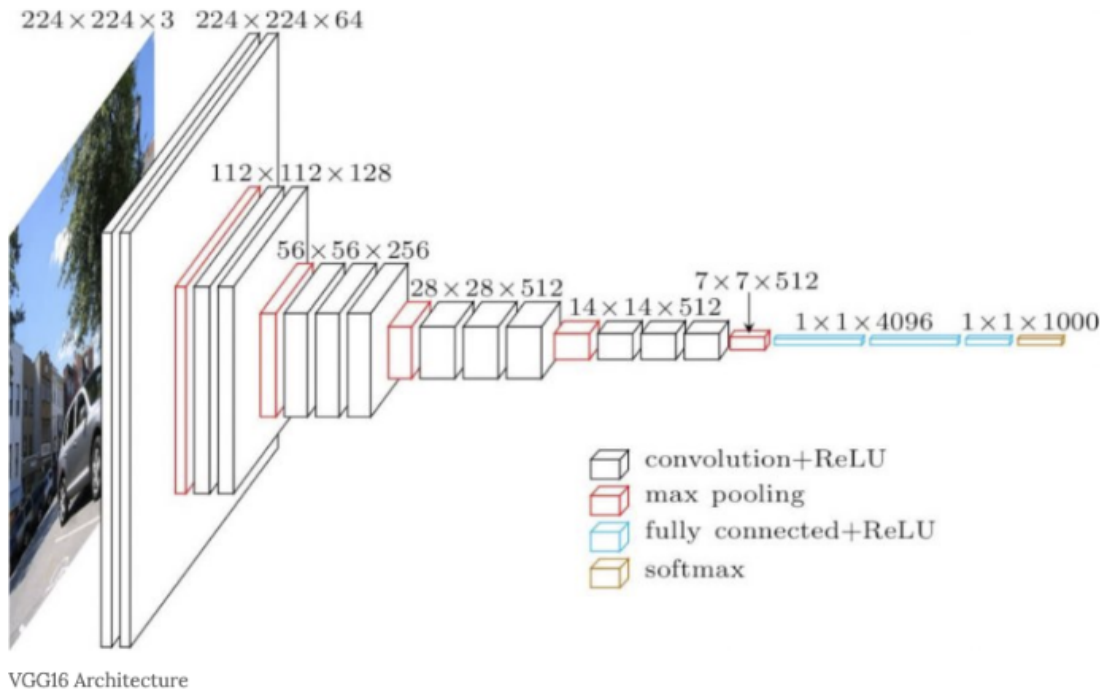
Table 3: ConvNet performance at a single test scale.

ConvNet config. (Table 1)	smallest image side		top-1 val. error (%)	top-5 val. error (%)
	train (S)	test (Q)		
A	256	256	29.6	10.4
A-LRN	256	256	29.7	10.5
B	256	256	28.7	9.9
C	256	256	28.1	9.4
	384	384	28.1	9.3
	[256;512]	384	27.3	8.8
D	256	256	27.0	8.8
	384	384	26.8	8.7
	[256;512]	384	25.6	8.1
E	256	256	27.3	9.0
	384	384	26.9	8.7
	[256;512]	384	25.5	8.0

5. Conclusion

Large-scale image classification을 위해서 19 weights layers CNN을 활용하였습니다. Representation depth는 Classification Accuracy에서 매우 유익하며, SOTA에 달성하는데 도움을 준다는 것을 알 수 있었습니다.

▼ VGG16 MODEL [Pytorch]



center). In one of the configurations we also utilise 1×1 convolution filters, which can be seen as a linear transformation of the input channels (followed by non-linearity). The convolution stride is fixed to 1 pixel; the spatial padding of conv. layer input is such that the spatial resolution is preserved after convolution, i.e. the padding is 1 pixel for 3×3 conv. layers. Spatial pooling is carried out by five max-pooling layers, which follow some of the conv. layers (not all the conv. layers are followed by max-pooling). Max-pooling is performed over a 2×2 pixel window, with stride 2.

```
import torch
import torch.nn as nn
import torch.nn.functional as F
from torchsummary import summary

device = "cuda:0" if torch.cuda.is_available() else "cpu"

VGG_16 = [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 'M', 512, 512, 512, 'M', 512, 512, 512, 'M']

class VGG16(nn.Module):
    def __init__(self, in_channels, n_classes):
        super(VGG16, self).__init__()
        self.in_channels = in_channels

        self.feature_extractor = self.create_conv_layers(VGG_16)

        self.classification = nn.Sequential(
            nn.Linear(512*7*7, 4096),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(4096, n_classes)
        )

    def forward(self, x):
        x = self.feature_extractor(x)
        x = x.view(-1, 512*7*7)
        x = self.classification(x)

        return x

    def create_conv_layers(self, architecture):
```

```

layers = []
in_channels = self.in_channels

for x in architecture:
    if type(x) == int:
        out_channels = x

        layers += [nn.Conv2d(in_channels = in_channels, out_channels = out_channels,
                               kernel_size = (3, 3), stride = (1, 1), padding = (1, 1)),
                    nn.BatchNorm2d(x),
                    nn.ReLU())
        in_channels = x
    elif x == 'M':
        layers += [nn.MaxPool2d(kernel_size = (2, 2), stride = (2, 2))]

return nn.Sequential(*layers)

model = VGG16(in_channels = 3, n_classes= 1000).to(device)
summary(model, input_size = (3, 224, 224))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
BatchNorm2d-2	[-1, 64, 224, 224]	128
ReLU-3	[-1, 64, 224, 224]	0
Conv2d-4	[-1, 64, 224, 224]	36,928
BatchNorm2d-5	[-1, 64, 224, 224]	128
ReLU-6	[-1, 64, 224, 224]	0
MaxPool2d-7	[-1, 64, 112, 112]	0
Conv2d-8	[-1, 128, 112, 112]	73,856
BatchNorm2d-9	[-1, 128, 112, 112]	256
ReLU-10	[-1, 128, 112, 112]	0
Conv2d-11	[-1, 128, 112, 112]	147,584
BatchNorm2d-12	[-1, 128, 112, 112]	256
ReLU-13	[-1, 128, 112, 112]	0
MaxPool2d-14	[-1, 128, 56, 56]	0
Conv2d-15	[-1, 256, 56, 56]	295,168
BatchNorm2d-16	[-1, 256, 56, 56]	512
ReLU-17	[-1, 256, 56, 56]	0
Conv2d-18	[-1, 256, 56, 56]	590,080
BatchNorm2d-19	[-1, 256, 56, 56]	512
ReLU-20	[-1, 256, 56, 56]	0
Conv2d-21	[-1, 256, 56, 56]	590,080
BatchNorm2d-22	[-1, 256, 56, 56]	512
ReLU-23	[-1, 256, 56, 56]	0
MaxPool2d-24	[-1, 256, 28, 28]	0
Conv2d-25	[-1, 512, 28, 28]	1,180,160
BatchNorm2d-26	[-1, 512, 28, 28]	1,024
ReLU-27	[-1, 512, 28, 28]	0
Conv2d-28	[-1, 512, 28, 28]	2,359,808
BatchNorm2d-29	[-1, 512, 28, 28]	1,024
ReLU-30	[-1, 512, 28, 28]	0
Conv2d-31	[-1, 512, 28, 28]	2,359,808
BatchNorm2d-32	[-1, 512, 28, 28]	1,024
ReLU-33	[-1, 512, 28, 28]	0
MaxPool2d-34	[-1, 512, 14, 14]	0
Conv2d-35	[-1, 512, 14, 14]	2,359,808
BatchNorm2d-36	[-1, 512, 14, 14]	1,024
ReLU-37	[-1, 512, 14, 14]	0
Conv2d-38	[-1, 512, 14, 14]	2,359,808
BatchNorm2d-39	[-1, 512, 14, 14]	1,024
ReLU-40	[-1, 512, 14, 14]	0
Conv2d-41	[-1, 512, 14, 14]	2,359,808
BatchNorm2d-42	[-1, 512, 14, 14]	1,024
ReLU-43	[-1, 512, 14, 14]	0
MaxPool2d-44	[-1, 512, 7, 7]	0
Linear-45	[-1, 4096]	102,764,544
ReLU-46	[-1, 4096]	0
Dropout-47	[-1, 4096]	0
Linear-48	[-1, 4096]	16,781,312
ReLU-49	[-1, 4096]	0
Dropout-50	[-1, 4096]	0
Linear-51	[-1, 1000]	4,097,000
Total params: 138,365,992		
Trainable params: 138,365,992		
Non-trainable params: 0		
Input size (MB): 0.57		
Forward/backward pass size (MB): 321.95		
Params size (MB): 527.82		
Estimated Total Size (MB): 850.35		

▼ Baseline [Pytorch]

```
import torch
import torch.nn as nn
from torch import optim
from torch.optim.lr_scheduler import StepLR

from torchvision import datasets
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import os
```

```

from torchvision import utils
import matplotlib.pyplot as plt
%matplotlib inline

import numpy as np
from datetime import datetime
import copy

```

```

path2data = 'data'
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Resize((224, 224)),
    transforms.Normalize(
        RGB_mean, RGB_std)
])
train_dataset = datasets.STL10(path2data, split='train',
                               download=True, transform=transform)
valid_dataset = datasets.STL10(path2data, split='test',
                                download=True, transform=transform)

```

```

meanRGB = [np.mean(x.numpy(), axis=(1, 2)) for x, _ in train_dataset]
stdRGB = [np.std(x.numpy(), axis=(1, 2)) for x, _ in train_dataset]

meanR = np.mean([m[0] for m in meanRGB])
meanG = np.mean([m[1] for m in meanRGB])
meanB = np.mean([m[2] for m in meanRGB])

stdR = np.mean([s[0] for s in stdRGB])
stdG = np.mean([s[1] for s in stdRGB])
stdB = np.mean([s[2] for s in stdRGB])

RGB_mean = [meanR, meanG, meanB]
RGB_std = [stdR, stdG, stdB]

```

```

train_dataloader = DataLoader(train_dataset, batch_size=4, shuffle=True)
valid_dataloader = DataLoader(valid_dataset, batch_size=4, shuffle=False)

```

```

import torch
import torch.nn as nn
import torch.nn.functional as F
# from torchsummary import summary

device = "cuda:0" if torch.cuda.is_available() else "cpu"

VGG_16 = [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 'M', 512, 512, 512, 'M', 512, 512, 512, 'M']

class VGG16(nn.Module):
    def __init__(self, in_channels, n_classes):
        super(VGG16, self).__init__()
        self.in_channels = in_channels

        self.feature_extractor = self.create_conv_layers(VGG_16)

        self.classification = nn.Sequential(
            nn.Linear(512*7*7, 4096),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(4096, n_classes)
        )

    def forward(self, x):

```

```

        x = self.feature_extractor(x)
        x = x.view(-1, 512*7*7)
        x = self.classification(x)

        return x

def create_conv_layers(self, architecture):
    layers = []
    in_channels = self.in_channels

    for x in architecture:
        if type(x) == int:
            out_channels = x

            layers += [nn.Conv2d(in_channels = in_channels, out_channels = out_channels,
                                kernel_size = (3, 3), stride = (1, 1), padding = (1, 1)),
                       nn.BatchNorm2d(x),
                       nn.ReLU())
            in_channels = x
        elif x == 'M':
            layers += [nn.MaxPool2d(kernel_size = (2, 2), stride = (2, 2))]

    return nn.Sequential(*layers)

model = VGG16(in_channels = 3, n_classes= 10).to(device)
# summary(model, input_size = (3, 224, 224))

```

```

def get_accuracy(model, data_loader, device):
    correct_pred = 0
    n = 0

    with torch.no_grad():
        model.eval()
        for X, y_true in data_loader:
            X, y_true = X.to(device), y_true.to(device)
            y_pred = model(X)

            _, predicted_labels = torch.max(y_pred, 1)

            n += y_true.size(0)
            correct_pred += (predicted_labels == y_true).sum()

    return correct_pred.float() / n

```

```

def train(train_loader, model, criterion, optimizer, device):
    model.train()
    running_loss = 0
    for X, y_true in train_loader:
        X, y_true = X.to(device), y_true.to(device)
        y_hat = model(X)
        loss = criterion(y_hat, y_true)
        running_loss += loss.item() * X.size(0)
        loss.backward()
        optimizer.step()

    epoch_loss = running_loss / len(train_loader.dataset)
    return model, optimizer, epoch_loss

```

```

def validation(valid_loader, model, criterion, device):
    model.eval()
    running_loss = 0
    for X, y_true in valid_loader:
        X, y_true = X.to(device), y_true.to(device)

```



```

y_hat= model(X)
loss = criterion(y_hat, y_true)
running_loss += loss.item() * X.size(0)

epoch_loss = running_loss / len(valid_loader.dataset)

return model, epoch_loss

```

```

def training_loop(model, criterion, optimizer, train_loader, valid_loader, epochs, device, print_every = 1):
    best_loss = 1e10
    train_losses = []
    valid_losses = []

    for epoch in range(epochs):
        model, optimizer, train_loss = train(train_loader, model, criterion, optimizer, device)
        train_losses.append(train_loss)

        with torch.no_grad():
            model, valid_loss = validation(valid_loader, model, criterion, device)
            valid_losses.append(valid_loss)

        if epoch % print_every == (print_every - 1):
            train_acc = get_accuracy(model, train_loader, device)
            valid_acc = get_accuracy(model, valid_loader, device)

            print(f'{datetime.now().time().replace(microsecond=0)} --- '
                  f'Epoch: {epoch}\t'
                  f'Train loss: {train_loss:.4f}\t'
                  f'Valid loss: {valid_loss:.4f}\t'
                  f'Train accuracy: {100 * train_acc:.2f}\t'
                  f'Valid accuracy: {100 * valid_acc:.2f}')

    return model, optimizer, (train_losses, valid_losses)

```

```

device = "cuda:0" if torch.cuda.is_available() else "cpu"

model = VGG16(3, 10).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr = 1e-1)

model, optimizer, _ = training_loop(model, criterion, optimizer,
                                     train_dataloader, valid_dataloader, 15, device)

```

Very Deep Convolutional Networks for Large-Scale Image Recognition

In this work we investigate the effect of the convolutional network depth on its accuracy in the large-scale image recognition setting. Our main contribution is a thorough evaluation of networks of increasing depth using an architecture with very small (3x3)

 <https://arxiv.org/abs/1409.1556>

